

Yoda Easy Guide

David Cohen

ECE Department – Carnegie Mellon University

Introduction

YODA, Yet another Ontological Dialog Architecture is a sophisticated tool to allow developers to quickly build intelligent spoken dialog systems. This tutorial will introduce the major components and capabilities of Yoda by walking through the creation of a simple messaging calendar assistant. We envision this system as a smartphone app, which the user can use to message their friends and coordinate meetings.

YODA's Dialog Functionality

The YODA dialog manager is designed to build dialog systems which combine the functions of information retrieval (IR), and information entry (IE). An IR dialog system's main purpose is to answer questions about a database of objects, such as restaurants, movies, or events in a calendar. A situated environment is also represented as a database, and IR dialog patterns still apply to asking questions about the environment.

An IE dialog system allows the user to describe objects and to have those descriptions used to modify the system's database, for example, adding new calendar events to a schedule. IE dialog patterns extend to command-and-control dialog, interactive grounding in partially-observed situated environments, crowdsourcing databases, and many other applications.

The following list shows the initial major categories of discourse unit supported by YODA:

- Information Retrieval
 - YN Question
 - WH Question
 - Search
- Information Entry
 - Command
 - Offer
 - Statement

A discourse unit is a dialog task, characterized by its DU-type and its semantic content. A discourse unit will stretch over multiple utterances, including initial presentations and related clarification, grounding dialog.

One of the major expected contributions of the YODA dialog manager is to support multiple discourse units gracefully. This will enable dialog systems which can maintain context across tasks, and support multiple simultaneous tasks.

YODA supports complex dialog systems by providing interfaces for specifying what types of discourse units are allowed for what types of database objects, as well as providing interfaces to non-dialog tasks which can be triggered by discourse units.

Our demonstration system can talk about people, meetings, times, and emails. Every meeting has a time and several attendees. Every email has a sender, a recipient, and some text content. We want the system to be able to set up a meeting, as well as answer basic queries about a user's schedule. We also want the system to be able to answer basic queries about their inbox, and to allow the user to dictate emails.

To keep the example simple, and to play to YODA's strengths, we assume that email dictation is handled by an external component, which the YODA dialog manager triggers as appropriate. The main reasons for this are: 1) a dialog system should not perform understanding on dictated text, so the processing of a dictated email is different than the processing of normal dialog; and 2) different ASR components will be appropriate for dictated email versus standard dialog.

SLU Input

The dialog manager takes as input an n-best list of hierarchical semantic parse results (defined by YODA's `SemanticsModel` class) along with corresponding confidence values. Our utterance representation is designed to allow natural utterances such as fragments and clarification utterances to be interpreted usefully, while still retaining principled semantic expressiveness.

We use an incomplete list of common semantic roles to denote the relationship between an action and the various entities under discussion. In the near term, these will likely include:

- Agent (often Subject in English syntax)
- Patient (often Direct Object in English syntax)
- Recipient (often Indirect Object in English syntax)
- FromLocation
- ToLocation
- Location
- FromTime
- ToTime
- Time

The utterance semantic representation also includes a dialog act, a polarity, and possibly other slots in the future. Some example utterances and corresponding semantic representations in JSON format are shown in Figure 1. Each role can either be unfilled, or filled with an object description.

Example a) shows a complete sentence and its interpretation. Notice that “want” is not used to determine the sentence’s action, rather it is interpreted as an indirect command. Similarly, “why don’t you X ?” should be interpreted as a command, rather than a WH question.

Example b) shows how two times can be represented within a single utterance.

Example c) shows the format that recognized named entities should be given in. The `nameGiven` slot within an object description is filled by the exact string used in the utterance. This string is then used to resolve the reference within later understanding modules. The same `nameGiven` slot is used for a named place, company, country, and so on.

The exact set of slots available to an object description correspond to the properties defined within the ontology. We also intend to provide an interface to support domain-specific higher-level understanding such as assessing fuzzy properties and resolving demonstratives for reference resolution. Details of this are outside the scope of this tutorial.

Build the Ontology

A YODA developer must then formalize the objects, actions and properties which are relevant to their domain by extending the YODA skeleton ontology with new classes, properties, values, and instances. Objects should be given appropriate slots, relations, and properties which correspond to the actual objects. There are several powerful tools available for editing OWL ontologies, but we recommend *protege*, and all our examples are shown in *protege*.

Figure 2 shows the class hierarchy for the domain ontology.

a) "I want to set up a meeting"

```
{ dialogAct: "command",  
  action: "Create",  
  patient: {  
    class: "Meeting"  
    number: "singular"  
    refType: "indefinite"  
  }  
}
```

b) "From two to three"

```
{ dialogAct: "fragment",  
  fromTime: {  
    class: "Time"  
    hour: "two"  
  }  
  toTime: {  
    class: "Time"  
    hour: "three"  
  }  
}
```

c) "With Sam"

```
{ dialogAct: "fragment",  
  adjunct: {  
    class: "Person"  
    nameGiven: "Sam"  
  }  
}
```

Figure 1: SLU input examples

Define the Dialog Functionality

Dialog Tasks

The developer defines what dialog tasks

Non-dialog Tasks

Several aspects of a non-dialog task will effect how the dialog manager interacts with them. In addition to implementing the execution of these tasks, the developer must define the information relevant for dialog management.

Figure 3 shows the definition for the `NonDialogTaskPreferences` class. It contains several parameters which determine the reward and penalty (negative reward) received by the system for executing a task, delaying execution of a task, and incorrectly executing a task or executing a task with incorrect parameters. These preferences are used by the dialog manager to perform decision-theoretic dialog management, and may be used to train reinforcement-learning dialog managers in the future.

There is another parameter, `alwaysRequireConfirmation`, that can be set to true if the dialog manager is required to get a confirmation from the user before performing an non-dialog task. Depending on the cost / benefit tradeoff, the dialog manager may decide to request a confirmation before executing a non-dialog task even if the flag is not set, and setting a high penalty for incorrect execution can make explicit confirmation requests more likely, but setting the flag allows the designer to enforce strict transparency.

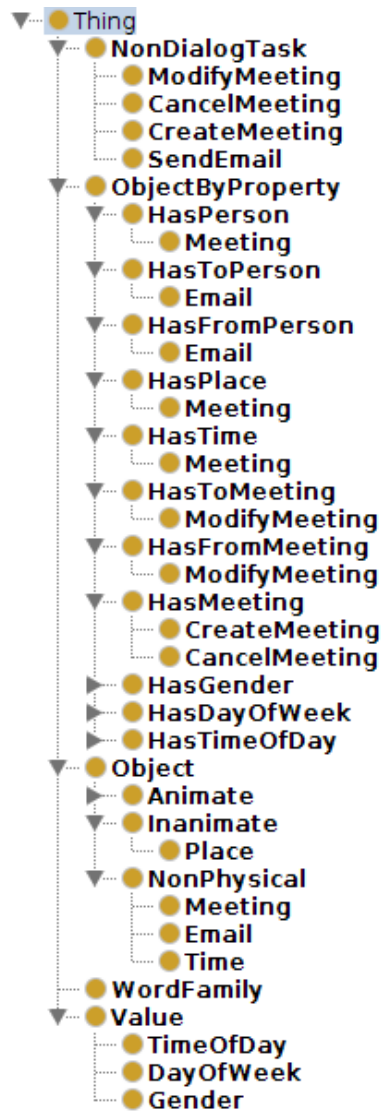


Figure 2: YODA sample domain ontology

In our Figure 5 line 5-7, we show the preferences declaration for the send-email task. Here, we set `alwaysRequireConfirmation` to true, so the user will always be asked for explicit confirmation before an email is sent, as in the following example:

```

User Send an email to Jerry.
System You want me to send an email to Jerry, right?
User Yes.
System Ok.
System <executes task>
  
```

In our Figure 6 line 5-10, we show the preferences declaration for the create-meeting task. Here, we set `alwaysRequireConfirmation` to false, so the system is allowed to perform the task without an explicit confirmation, as in the following example:

```

User Set up a meeting this afternoon with Jerry in my office.
System Ok.
System <executes task>
  
```

```

1 public class NonDialogTaskPreferences {
2     // if explicit confirmation is always required before the
3     // task is executed, set alwaysRequireConfirmation to true
4     public boolean alwaysRequireConfirmation;
5
6     // all the rewards and penalties are given as positive
7     // numbers, penalties are converted to negative numbers for
8     // decision-making
9     public double penaltyForDelay;
10    public double rewardForCorrectExecution;
11    public double penaltyForIncorrectExecution;
12
13    // define the slots that are required to be defined by the
14    // user before execution is possible
15    public Set<String> slotsRequired;
16
17    public NonDialogTaskPreferences(
18        boolean alwaysRequireConfirmation,
19        double penaltyForDelay,
20        double rewardForCorrectExecution,
21        double penaltyForIncorrectExecution,
22        Set<String> slotsRequired) {
23        this.alwaysRequireConfirmation =
24            alwaysRequireConfirmation;
25        this.penaltyForDelay = penaltyForDelay;
26        this.rewardForCorrectExecution =
27            rewardForCorrectExecution;
28        this.penaltyForIncorrectExecution =
29            penaltyForIncorrectExecution;
30        this.slotsRequired = slotsRequired;
31    }
32 }

```

Figure 3: The NonDialogTask interface

The last parameter of the NonDialogTaskPreferences class sets requirements for what must be specified for a task to be executable. As is shown in the preference declarations from Figures 5 and 6, sending an email only requires a sender, while creating a meeting requires a description of a meeting with at least one time, place, and attendee. The dialog manager will not consider executing a task until all its required slots are filled, and will instead engage in slot-filling dialog, as in the following example:

```

User Set up a meeting this afternoon.
System Where will this meeting be?
User In my office.
System Who is this meeting with?
User Jerry.
System Ok.
System <executes task>

```

After basic preferences have been defined, the NonDialogTask interface, shown in figure 4, must be implemented for each non-dialog task. Two implementations of this are shown in this tutorial, SendEmailTask is shown in figure 5, and CreateMeetingTask is shown in figure 6

The assessExecutability method is a task-specific method for informing the dialog manager how likely it is a task will succeed. For example, sending an email may be

```

1 public interface NonDialogTask {
2
3     public enum TaskStatus {SUCCESSFULLY_COMPLETED,
4         CURRENTLY_EXECUTING_BLOCKING,
5         CURRENTLY_EXECUTING_NOT_BLOCKING, FAILED}
6
7     // return preferences object
8     public NonDialogTaskPreferences getPreferences();
9
10    // the probability that the taskSpec can be executed (0-1)
11    public double assessExecutability(SemanticsModel taskSpec);
12
13    // return the string ID of the executing task (taskID)
14    public String execute(SemanticsModel taskSpec);
15
16    // return the string status indicator for the taskID
17    public TaskStatus status(String taskID);
18
19 }

```

Figure 4: The NonDialogTask interface

guaranteed to fail if the device does not have working internet. For situated agents such as robots, there is always some non-negligible chance that an attempted physical action will fail. The dialog manager will incorporate these probabilities into its decision-making process. Intuitively, it will hesitate to perform actions which are likely to fail, and eagerly execute actions which are likely to succeed.

The execute method executes the actual task. This will call developer-defined functions and libraries to actually send emails, reserve meeting rooms, etc.

The status method checks a task's status, which can be one of the following:

- SUCCESSFULLY_COMPLETED
- CURRENTLY_EXECUTING_BLOCKING
- CURRENTLY_EXECUTING_NOT_BLOCKING
- FAILED

When a task is blocking, the dialog system is temporarily de-activated. In this tutorial, we assume that a separate module is called for email dication, so the send-email task is blocking.

```

1 public class SendEmailTask implements NonDialogTask {
2     private static Integer instanceCounter = 0;
3     private static Map<String, TaskStatus> executionStatus =
4         new HashMap<>();
5     private static NonDialogTaskPreferences preferences =
6         new NonDialogTaskPreferences(true, 1, 20, 20,
7             new HashSet<>(Arrays.asList("hasToPerson")));
8
9     @Override
10    public NonDialogTaskPreferences getPreferences() {
11        return preferences;
12    }
13
14    @Override
15    public double assessExecutability(SemanticsModel taskSpec) {
16        // this works because assessExecutability is never
17        // called unless all the required slots are present
18        return 1.0;
19    }
20
21    @Override
22    public String execute(SemanticsModel taskSpec) {
23        System.out.println("Executing task: send email");
24        System.out.println(taskSpec);
25        String ans = "SendEmailTask:" + instanceCounter.
26            toString();
27        instanceCounter += 1;
28        // There is no real implementation, so we just set
29        // the status to successfully completed
30        executionStatus.put(ans,
31            TaskStatus.SUCCESSFULLY_COMPLETED);
32
33        /*
34         * In the envisioned implementation, a separate
35         * dictation program would run while the other
36         * program is run, the dialog manager should block, as
37         * in the following commented lines
38         */
39        // executionStatus.put(ans,
40        //     TaskStatus.CURRENTLY_EXECUTING_BLOCKING);
41        // dictationProgram.run();
42
43        System.out.println("taskID: " + ans);
44        return ans;
45    }
46
47    @Override
48    public TaskStatus status(String taskID) {
49        return executionStatus.get(taskID);
50    }
51 }

```

Figure 5: The SendEmailTask implementation

```

1 public class CreateMeetingTask implements NonDialogTask {
2     private static Integer instanceCounter = 0;
3     private static Map<String, TaskStatus> executionStatus =
4         new HashMap<>();
5     private static NonDialogTaskPreferences preferences =
6         new NonDialogTaskPreferences(false, 1, 20, 15,
7             new HashSet<>(Arrays.asList(
8                 "hasMeeting.hasTime",
9                 "hasMeeting.hasPlace",
10                "hasMeeting.hasPerson"))));
11
12     @Override
13     public NonDialogTaskPreferences getPreferences() {
14         return preferences;
15     }
16
17     @Override
18     public double assessExecutability(SemanticsModel taskSpec) {
19         return 1.0;
20     }
21
22     @Override
23     public String execute(SemanticsModel taskSpec) {
24         System.out.println("Executing task: create meeting");
25         System.out.println(taskSpec);
26         String ans = "CreateMeetingTask:" + instanceCounter.
27             toString();
28         instanceCounter += 1;
29         // There is no real implementation, so we just set the
30         // status to successfully completed
31         executionStatus.put(ans,
32             TaskStatus.SUCCESSFULLY_COMPLETED);
33         System.out.println("taskID: " + ans);
34         return ans;
35     }
36
37     @Override
38     public TaskStatus status(String taskID) {
39         return executionStatus.get(taskID);
40     }
41 }

```

Figure 6: The SendEmailTask implementation